

Estimation Principles for Structural Models

Janoš Gabler

July 3, 2019

Outline

- 1 Introduction
- 2 Basic Ideas
- 3 Numerical Optimization
- 4 Monte Carlo Results
- 5 Some Theory
- 6 Estimagic

Introduction

Introduction

- ▶ Structural models are used for ex-ante policy evaluation
- ▶ They are tailored to one policy question
- ▶ Their parameters have to be estimated
- ▶ We fully abstracted from the estimation problem
- ▶ This time we fully abstract from complex models
 - ▶ Only estimate means and linear models

Goals for this lecture

- ▶ Explain the core estimation principles on toy models
 - ▶ Maximum Likelihood
 - ▶ Method of Simulated Moments
- ▶ Make you aware of typical numerical problems
- ▶ Introduce you to numerical optimization
- ▶ Explain why it is important to separate models from estimators

Basic Ideas

Preparation

Parametric Models

- ▶ A model is called parametric if all functions and distributions are specified up to a finite set of parameters
- ▶ Given a parameter vector and a parametric model:
 - ▶ A dataset can be simulated
 - ▶ Will be used for Method of Simulated Moments
 - ▶ The probability of observing a certain data point can be calculated
 - ▶ Will be used for Maximum Likelihood Estimation

Examples for Parametric and Non-Parametric Models

Example	Assumptions
$y_i = m(x_i, u_i)$	Smoothness of m , iid sampling
$y_i = m(x_i) + u_i$	+ additivity of errors
$y_i = m(x_i' \beta) + u_i$	+ single linear index
$y_i = x_i \beta + u_i$	+ m is the identity function
$u_i \sim \mathcal{N}(\mu, \sigma^2)$	+ distributional assumption

Example 1: Mean of a Normal Distribution

- ▶ Dataset: iid sample of variable y_i
- ▶ Model: $y_i \sim \mathcal{N}(\mu, \sigma^2)$
- ▶ Parameters to estimate: μ, σ

Maximum Likelihood


Example 1

- ▶ The model implies:

$$f(y_i) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\left(-\frac{1}{2\sigma^2}(y_i-\mu)^2\right)} \quad (1)$$

- ▶ f is the probability density function (pdf) of y_i
- ▶ Interpretation: y_i varies, μ and σ are fixed
- ▶ Density of whole sample is just product of individual densities

Graphical Intuition



Basic Idea of Maximum Likelihood

$$f(y_i) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\left(-\frac{1}{2\sigma^2}(y_i-\mu)^2\right)} = \phi(y_i, \mu, \sigma) = l_i(\mu, \text{sigma}) \quad (2)$$

- ▶ l is the likelihood contribution of individual i
- ▶ Interpretation: y_i is fixed, μ and σ vary
- ▶ Likelihood of whole sample is just product of individual likelihoods
- ▶ Use the parameters that maximize the likelihood of the sample as estimates

Python Implementation

```
import numpy as np
import scipy

def likelihood(mu, sigma, sample):
    likelihoods = scipy.stats.norm.pdf(sample, loc=mu, scale=sigma)
    return np.prod(likelihoods)
```

Notebook

Look at `likelihood_example_1.ipynb`

Some Insights from the Example

- ▶ Use log-likelihood to avoid numerical problems
- ▶ Likelihood estimates can be counter-intuitive
- ▶ Larger sample \rightarrow more curved likelihood \rightarrow more precision
 - ▶ Standard errors will depend on curvature of likelihood
- ▶ Mean is estimated quite precisely, even in small samples

A Note on Terminology

- ▶ Parameters are constants, not random variables
- ▶ Maximum likelihood estimates are not “the most likely parameters”
- ▶ They are the parameter values that make the observed sample most likely!

Method of Simulated Moments

Basic Idea

- ▶ Remember: If we have a fully parametric model and a parameter vector, we can simulate a dataset generated by the model with that parameter vector
- ▶ For the true parameter vector, simulated and observed sample should be similar
- ▶ MSM: Take the parameter vector that produces the dataset which is most similar to the empirical data as estimate

What Does Similar Mean?

- ▶ Depends on the model
- ▶ Typically: Key moments are similar
- ▶ Selection of key moments depends on model
- ▶ Requires weighting matrix

Example 1: Which moments

- ▶ Key moments: unconditional mean and variance
- ▶ Looks like chicken-or-egg problem
- ▶ Only an artifact of very simple example

Example 1: Objective function

- ▶ empirical sample: $Y = y_1, y_2, \dots, y_n$
- ▶ parameter vector: $\theta = (\mu, \sigma)^T$
- ▶ simulated sample, given θ : $\hat{Y}^\theta = \hat{y}_1^\theta, \hat{y}_2^\theta, \dots, \hat{y}_n^\theta$
- ▶ empirical moments: $m^{emp} = (\text{mean}(Y), \text{var}(Y))^T$
- ▶ simulated moments: $m(\theta) = (\text{mean}(\hat{Y}^\theta), \text{var}(\hat{Y}^\theta))^T$
- ▶ Objective: $C(\theta) = (m(\theta) - m^{emp})^T W (m(\theta) - m^{emp})$
- ▶ θ^{MSM} minimizes this function

Python Implementation

```
import numpy as np
import scipy

def msm_criterion(mu, sigma, sample):
    simulated_sample = np.random.normal(loc=mu, scale=sigma, size=len(sample))
    m_emp = np.array([[sample.mean()], [sample.std()]])
    m_sim = np.array([[simulated_sample.mean()], [simulated_sample.std()]])
    diff = m_emp - m_sim
    w = np.eye(2)
    return diff.T.dot(w).dot(diff)[0][0]
```


Notebook

Look at `msm_example_1.ipynb`

Some Insights from the Example

- ▶ Precision is now limited by two factors:
 - ▶ Sample variation
 - ▶ Simulation noise
- ▶ Likelihood needs the same assumption and will always be more efficient!

Why Would We Ever Use MSM?

- ▶ Likelihood often involves high dimensional integrals
- ▶ Likelihood is harder to derive
- ▶ Some say, MSM is more transparent
- ▶ Data might be on a slow server
- ▶ Different datasets

Numerical Optimization

Why Numerical Optimization

- ▶ Typically optimum cannot be calculated in closed form
 - ▶ Derivatives are too complicated to evaluate
 - ▶ Non differentiable functions
 - ▶ First order conditions are too complicated
- ▶ Numerical optimization works in these cases but:
 - ▶ Typically no guarantee of global optimum
 - ▶ Can be slow
 - ▶ Less precise than closed form
- ▶ WLOG, from now on only talk about minimization

We Already did it

- ▶ For surface plots, we evaluated objective over whole reasonable parameter space
- ▶ Know coordinates of optimum
- ▶ This methods is called grid search or brute force
- ▶ Infeasible for high dimensional parameter vectors:
 - ▶ Grid with 100 points in each dimension
 - ▶ 50 parameters
 - Criterion has to be evaluated at 10^{100} gridpoints
- ▶ Need smarter algorithms to save evaluations!

Gradient Based vs. Gradient Free

- ▶ Gradient based algorithms
 - ▶ Require differentiable objective functions
 - ▶ Use first derivative to determine direction of step
 - ▶ Use second derivative to determine length of step
 - ▶ Need much fewer function evaluations
- ▶ Gradient free algorithms
 - ▶ Work with non-smooth functions
 - ▶ Basically a trial-and-error approach

Termination Criteria

- ▶ ftol
- ▶ gtol
- ▶ xtol

Local vs. Global Optimizers

- ▶ Local Optimizers:
 - ▶ Start from start parameters and go down-hill
 - ▶ Stop in first local minimum
- ▶ Genetic Algorithms (typically global):
 - ▶ Need bounds on parameter space
 - ▶ Sample uniformly from parameters space to get start population
 - ▶ Evolve to next generation by:
 - ▶ Killing worst parameters
 - ▶ Produce offspring of good parameters
- ▶ Pseudo global optimizers:
 - ▶ Do local optimization from random start values

Notebook

look at 1d_minimization.ipynb

Insights from the Example

- ▶ Sampling scheme can be very important
- ▶ Robust optimizers can solve some problems, but maybe you don't need them!
- ▶ Always understand the problem you want to solve!

Smooth Python Implementation

```
import numpy as np
import scipy

def smooth_msm_criterion(mu, sigma, sample):
    np.random.seed(5471)
    simulated_sample = np.random.normal(size=len(sample)) * sigma + mu
    m_emp = np.array([[sample.mean()], [sample.std()]])
    m_sim = np.array([[simulated_sample.mean()], [simulated_sample.std()]])
    diff = m_emp - m_sim
    w = np.eye(2)
    return diff.T.dot(w).dot(diff)[0][0]
```

Monte Carlo Results

Notebook

look at monte_carlo.ipynb

Some Theory

Closed Form Solution for Likelihood Function in Example 1

$$L(\mu, \sigma) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{1}{2\sigma^2}(y_i - \mu)^2\right)} \quad \text{Likelihood}$$

$$\ell(\mu, \sigma) = \sum_{i=1}^n -\frac{1}{2} \ln(2\pi\sigma) - \frac{1}{2\sigma^2}(y_i - \mu)^2 \quad \text{Log-Likelihood}$$

$$\frac{\partial \ell}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mu) \stackrel{!}{=} 0 \quad \text{FOC}$$

$$\mu^* = \frac{1}{n} \sum_{i=1}^n y_i$$

Example 2: Linear Model

- ▶ Dataset: iid sample of variables y_i, x_i
- ▶ Model: $y_i = x_i' \beta + u_i, u_i \sim \mathcal{N}(0, \sigma^2)$
- ▶ Parameters to estimate: β, σ

Example 2: Linear Model

- ▶ Dataset: iid sample of variables y_i, x_i
- ▶ Model: $y_i = x_i' \beta + u_i, u_i \sim \mathcal{N}(0, \sigma^2)$
- ▶ Parameters to estimate: β, σ
- ▶ You might be tempted to call this an OLS model
- ▶ Later you'll see why this is not a good idea

Likelihood Function of Linear Model

$$f(y_i | \beta, \sigma, x_i) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\left(-\frac{1}{2\sigma^2} (y_i - x_i' \beta)^2\right)}$$
$$L(\beta, \sigma) = \prod_{i=1}^n f(y_i | \beta, \sigma, x_i)$$

- ▶ Take logs
- ▶ Set FOCs to zero
- ▶ If you make no mistake, you will see that β^* equals the OLS estimator

The Deep Reason (simplified!)

- ▶ It can be shown that:
 - ▶ Maximum Likelihood is asymptotically efficient
 - ▶ No unbiased estimator has smaller variance than maximum likelihood (Cramér-Rao Bound)
- ▶ From Gauss Markov Theorem we know:
 - ▶ OLS estimator is Best Linear Unbiased Estimator
 - ▶ Remember: sample mean = OLS on constant
- ▶ Since OLS and Maximum Likelihood estimators are both optimal, they have to coincide



MSM of Linear Model

- ▶ Identifying assumption: u_i is uncorrelated with x_i
- ▶ $E(x_i u_i) = E(x_i (y_i - x_i' \beta)) = 0$
- ▶ Gives us k moment conditions for k parameters
- ▶ Those are the moments used when Linear model is estimated with GMM
- ▶ Could use the simulated version of those moments to estimate Linear model with MSM
- ▶ Reminder: This is just to explain the method

Estimagic

Estimagic

- ▶ Wrap local, global and pseudo global optimizers from Pygmo, Scipy and TAO
- ▶ Standard errors for likelihood and MSM models
- ▶ Elegant interface for typical constraints
 - ▶ e.g. estimate covariance matrices, probabilities, ...
- ▶ Interactive dashboard with live convergence plots



Some Tipps

- ▶ Start by implementing toy examples
- ▶ When you hit problem, reproduce and solve them in minimal examples
- ▶ Use as much existing code as possible
- ▶ If you think you need a fancy optimizer, you probably don't
- ▶ If you ever do structural econometrics: use estimagic!
- ▶ But maybe wait a few months . . .